



Pokémon TCG Policy Network

64 state tokens | d = 384 | 5 layers | 6 attention heads per layer


State processing (game state)



Option processing (candidate actions)



Training / evaluation branches · omitted from NumPy policy inference



Card / attack representation = learned ID embedding + projected static features.